

A Complete Guide to Building a Beowulf Cluster

Luis Pimentel

VERSION 1.0, JUNE 2016, ROCKDALE CAREER ACADEMY

Contents

1	Introduction	1
2	Fundamentals	2
2.1	High Performance Computing, Computer Clusters and Parallel Processing	2
2.2	Beowulf Clusters	2
3	Getting Started	3
3.1	Hardware	3
3.2	Software	3
3.3	How to Use This Guide	4
4	Building the Cluster	5
5	Head Node Configuration	6
5.1	Configuring Hosts File	6
5.2	Configuring IP Addresses	7
5.2.1	Configuring IP Addresses Method 1	7
5.2.2	Configuring IP Addresses Method 2	7
6	Slave Node Configuration	9
6.1	Hosts File	9
6.2	IP Addresses	9
7	Setting Up the Network File System	10
7.1	Installing NFS on Head Node	10
7.2	Installing and Mounting NFS on Slave Nodes	10
8	Setting Up SSH Password-less Communication	12
9	Installing Compilers	14
10	Installing MPICH2	15
10.1	Getting Examples	15
11	Testing and Benchmarks	16
12	Linux References	17
13	Other References	18
14	Appendix	19

1 Introduction

2 Fundamentals

The purpose for this section is to provide an overview of the different technologies that this project entails. This will begin with an explanation of High Performance Computing, computer clusters, and parallel processing. Afterwards is an explanation of what exactly a Beowulf Cluster is and why we chose to build this type of cluster.

The purpose of this section is not to go completely in depth with an explanation of these technologies, but to provide a basic preview that can be useful for beginners. With that being said I highly encourage further reading on each of these topics and other related topics.

2.1 High Performance Computing, Computer Clusters and Parallel Processing

In simple terms, High Performance Computing (HPC) is the use of parallel processing to run advanced programs. These programs are run on HPC systems so that they can be processed much faster and efficiently. HPC systems are often used by advanced researchers in the scientific field and in government. Typically any organization that is in need of powerful computing is likely to have an HPC system, whether it is made up of special custom components or everyday computer hardware.

One way to take advantage of HPC is through the use of computer clusters. A computer cluster is typically a group of computers, in which each computer in the cluster is called a *node*. Each node has its own individual resources, such as a CPU, memory, and an Operating System. These nodes are typically connected through a Local Area Network in order to communicate with each other. Clusters have become relatively easier to build over time, especially through the use of off-the-shelf computers and Linux operating systems. They can be very low-budget allowing practically anyone with spare computer hardware to build them.

These clusters use parallel processing to behave like "super computers". Parallel processing splits a task into multiple parts, giving each individual node a part of the task. This makes use of multiple CPUs at the same time, causing the program to be processed in less time than on a typical computer.

Through computer clusters and their use of parallel processing, we essentially create a "super computer". There are many different applications for these type of computers. Such applications have included protein folding simulation, climate prediction and graphics rendering. As said before, practically every industry in need of the processing power that computer clusters have to offer, is likely to own or use one of these machines.

There are many different types of clusters out there, classified by the computer architecture used (meaning how it is built and what hardware/software is being used). We decided to build a Beowulf Cluster.

2.2 Beowulf Clusters

Beowulf Clusters differ from other clusters in that any old computer capable of running Linux and capable of networking can be used to create a cluster. Furthermore each individual computer can be different from the other. You can use different computers to make up your cluster, making it an inexpensive choice for High Performance Computing. A typical Beowulf Cluster consists of one head node and several other slave nodes. These slave nodes can be added as needed, increasing the power of the cluster. The head node often acts as a central server, controlling all of the slave nodes. Other advantages to this type of cluster is that standards and tools have been developed, making it easier to develop programs for these machines.

For its inexpensiveness and relative easiness, we decided that building a Beowulf Cluster would be the best approach to building our very own "super computer". I used many of the old and new spare computers that we had available in our lab, and even build up a few of them.

3 Getting Started

3.1 Hardware

To build the Beowulf Cluster you are going to need a few computers. Technically you only need a minimum of 2 computers: one to act as the head node and another to act as the slave node. In our computer science room we used one head node and 8 slave computers. Our head node was a and our slave nodes consisted of .

When thinking about the number of nodes you are going to have, it is important to consider various factors including how much computing power you need, how much you are willing to spend and whether or not you have the infrastructure necessary to build a larger cluster. While relatively low compared to older "super computers", computer clusters still typically consume a lot of power. Your ability to store and network these computers should also be considered.

Each node in the cluster has a few hardware requirements. Each node will need a CPU, RAM, graphics, networking and storage. While you do not need the most cutting edge hardware, the better your processor is, the more powerful your computer cluster will be. As said before each of the individual nodes do not have to be the same. You can mix and match different motherboards, processors, etc. The only requirements are a basic Input/Output system and the ability to run Linux.

Because computer clusters communicate with each other through the use of a Local Area Network, basic networking equipment will be needed. The only networking requirements consist of a networking switch with enough ports to support the amount of nodes. Enough Ethernet CAT cable for each node is also required.

An Internet connection will be needed to download key software for this project.

3.2 Software

For this cluster we used Ubuntu Server as our operating system. While some experience with Linux is recommended, it is not fully required. This document will include basic tutorials for the necessary Linux commands.

[You can download the latest version of Ubuntu Server here](#)

Furthermore for our project I chose to not use a Graphics User Interface (GUI), meaning I used the default Command Line Interface (CLI) that comes after installing Ubuntu Server. I did this so that I could further add to my experience learning Linux (even if it was the hard way). Overtime I also found it more efficient. While you can opt to use a GUI rather than just the CLI, many of the commands necessary to run Linux are run in a terminal. However I do recommend just using the CLI, as it is a much more fun learning experience.

The rest of the software used in this project consists of packages that will be installed and explained further along the project.

While this project will provide a brief explanation of many of the commands used in this project, it is recommended that you do further reading to Linux operating systems.

Basic networking knowledge is also needed for this project. This document will not go into detail with some of these networking topics, so further reading into networking topics such as IP addressing is advised.

3.3 How to Use This Guide

I designed this guide to be used on chronological order. You should complete the steps as they are laid out.

I began with instructions for configuring the head node and then move on to configuring the slave nodes. However, as you progress through the guide, you will have to switch between working on the head node and working on the slave nodes. It is of extreme importance that you know which node you are supposed to be working on, as some commands are only supposed to be run on either machine.

As said before this guide provides not only instructions, but explanations for the different technologies you are using and how they are essential to the design of this cluster. Throughout the guide I will provide links for further reading on different topics.

Explanations for Linux procedures specific to a step in the project will be explained in that section of the project. Other general Linux explanations will be presented towards the end of the document in its own section. Feel free to reference this section throughout the project.

Other useful references will also be provided towards the end of the document.

4 Building the Cluster

The following are general instructions for setting up you computer cluster hardware. This part can vary depending on the setup that you want.

Instructions

- a. Begin by setting up you Head node. This should be your most powerful computer available. It is recommended to place this node separate from the rest of the nodes in the cluster. You should be able to easily distinguish this node from the others. This node should also have its very own Monitor, keyboard and mouse setup.
- b. Begin setting up the Slave nodes. Typically these nodes should be grouped together and labeled. The design of the cluster is completely up to you. You can have the computers stacked up on each other, or placed side by side on a rack.

Things to consider in this step include distance to the central networking switch and distance to a power source. Feel free to use power strips. Be careful with power however, as you do not want to overload one power outlet with too many computers.

- c. Proceed by setting up a central networking switch. When choosing the location of the switch, keep in mind the distance of the nodes.
- d. Begin networking the cluster. Connect an Ethernet CAT cable to each of the nodes, and connect each cable to a port on the central networking switch.

Consider cable management in this step as you do not want to create a tangle of wires. If possible, try to color code these cables or label them to easily identify them when troubleshooting networking issues. It is good practice to place each cable into the ports on the switch by order of node i.e. node 1 to port 1...etc.

- e. Go ahead and plug power strips to each of the nodes. However you should not turn on any of the nodes yet, as they do not have an operating system installed.

After completing these steps you should have successfully built the infrastructure necessary to create a Beowulf Cluster. We now proceed to configuring each of the nodes with the necessary software.

5 Head Node Configuration

Configuring the Head node is the first step in building the Beowulf cluster. This is the most important node in the cluster as we will run all of the jobs. This node will also act as a central server, providing many of the services described later in this document.

As said before this should be your most powerful machine, it should be in a distinguished area, with its own monitor, keyboard and mouse.

Many of the naming conventions that we used in this project can be changed depending on your preference. This includes usernames, hostnames, IP addresses etc. Feel free to change this for your own project. However make sure to document which things you are changing so that you do not cause confusion in future steps.

Instructions

- a. After having downloaded the latest version of Ubuntu Server, install this to the head node.
For more information on how to install Ubuntu reference the "Other References" section of this document.
- b. You should set the hostname of this computer to headwolf.
- c. You will need to create a user with the name of mpiuser. Set a password for this user.

5.1 Configuring Hosts File

The hosts file is used so that nodes can communicate with each other without the need to remember and type out each individual node's IP address of which you are trying to communicate with. It maps out each node's IP address to each node's hostname.

You will manually have to configure the hosts file on every machine.

Instructions

- a. Open up the hosts file within the /etc directory

```
sudo nano /etc/hosts
```

- b. Configure it to look like the following

```
127.0.0.1      localhost
10.0.0.100     headwolf
10.0.0.101     wolf01
10.0.0.102     wolf02
10.0.0.103     wolf03
```

- c. Add as many nodes as you have set up.
- d. Save file and exit out.

The 127.0.0.1 to localhost line is linking localhost to the computer's loopback address. This is necessary for network troubleshooting and does not pertain to building the Beowulf Cluster.

5.2 Configuring IP Addresses

To setup our cluster's network we need to configure the IP addresses of each node. We need to establish static IP addresses on each node that match the IP addresses mapped on the hosts file.

Because for this project we are using Ubuntu Server's default CLI we will have to do this by editing a file that sets the IP addresses to each Network Interface Card (NIC). The configuration of this file will be different depending on your setup.

If you have access to your Internet network and can configure your router to work with the IP addresses set to the nodes then proceed to step 5.2.1

If your Internet connection uses DHCP and you do not have access to your network refer to section 5.2.2

5.2.1 Configuring IP Addresses Method 1

Instructions

- a. Identify your current IP configuration.

```
ifconfig
```

- b. Your output should look like the following. Make note of the name used to identify your network card.

- c. To configure a static IP address to be used, edit the following file

```
sudo nano /etc/network/interfaces
```

To look like the following.

```
auto eth0
iface eth0 inet static
address 10.0.0.100
gateway 10.0.0.1
netmask 255.255.255.0
network 10.0.0.0
broadcast 10.0.0.255
```

- d. Restart your networking service for these settings to take place

```
sudo /etc/init.d/networking restart
```

**** Your gateway address should be the address of your router

**** Replace eth0 with the name of your NIC found in step 1. If it is the same do not replace.

5.2.2 Configuring IP Addresses Method 2

The following method should be used if your Internet connection uses DHCP and you only have 1 NIC in each of your nodes. If this is the case for you, you will have to switch back and forth between a Static address and a DHCP address for your NIC. This method is useful for those who do not access to their network (for example in a school environment). Alternatively if you do not have configuration access to the Internet network you are using, you can install your own router in between, and set that router up to work with the IP address of your cluster when an Internet

connection is needed. Then you can use method 1 to configure the IPs of your cluster. **This is recommended.**

Instructions

- a. Edit the following file

```
sudo nano /etc/network/interfaces
```

To look like the following.

```
auto eth0
iface eth0 inet static
address 10.0.0.100
gateway 10.0.0.1
netmask 255.255.255.0
network 10.0.0.0
broadcast 10.0.0.255

auto eth0
iface eth0 inet dhcp
```

These are the two configurations that will be used by your network card.

- b. When using either the static IP (for the cluster) or the DHCP address (for Internet), you must comment out the one you are not using. For example if you were using DHCP, your settings would look like the following.

```
#auto eth0
#iface eth0 inet static
#address 10.0.0.100
#gateway 10.0.0.1
#netmask 255.255.255.0
#network 10.0.0.0
#broadcast 10.0.0.255

auto eth0
iface eth0 inet dhcp
```

While if you were using the Static IP for your cluster

```
auto eth0
iface eth0 inet static
address 10.0.0.100
gateway 10.0.0.1
netmask 255.255.255.0
network 10.0.0.0
broadcast 10.0.0.255

#auto eth0
#iface eth0 inet dhcp
```

- c. Always remember to restart your networking settings.

```
sudo /etc/init.d/networking restart
```

******I had issues with the formatting of the document. The part of the configuration file that is commented should not be in italics for you. Simply add a (hashtag) to the lines you wish to comment out.**

6 Slave Node Configuration

6.1 Hosts File

To set up the hosts file for the slave nodes, use the same steps used to set up the hosts file on the Head node. (Section 5.1)

Repeat these steps for each Slave nodes in your cluster.

6.2 IP Addresses

To set up the IP addressess of the Slave nodes in your cluster, use the same steps used for the Head node. (Section 5.2)

The Static IP set in the `/etc/network/interfaces` configuration file should follow the following format:

```
auto eth0
iface eth0 inet static
address 10.0.0.XXX
gateway 10.0.0.1
netmask 255.255.255.0
network 10.0.0.0
broadcast 10.0.0.255
```

For example if you were configuring the 6th slave node in your cluster:

```
auto eth0
iface eth0 inet static
address 10.0.0.106
gateway 10.0.0.1
netmask 255.255.255.0
network 10.0.0.0
broadcast 10.0.0.255
```

Always remember to restart your networking settings.

```
sudo /etc/init.d/networking restart
```

7 Setting Up the Network File System

Network File System is a computer protocol that allows networked computers to access each other's file systems as if they were stored locally. We will use NFS to share the Head node's `/home/mpiuser` directory with the other nodes. Because every node has the username **mpiuser**, each node already has this directory. Therefore the `/home/mpiuser` directory is shared by the whole cluster. This is necessary because the nodes need to be able to access the jobs that we are running. Having the nodes virtually share the same directory also comes in handy as we install SSH for communication later on.

7.1 Installing NFS on Head Node

NFS must first be installed on the Head node from which we will share this directory.

To install NFS service on the Head node, run the following command:

```
sudo apt-get install nfs-kernel-server
```

Next export the `/home/mpiuser` directory so that the other nodes can have access to it. Do this by editing the `/etc/exports` file:

```
sudo nano /etc/exports
```

And add the following line to the end of the file:

```
/home/mpiuser *(rw,sync,no_subtree_check)
```

Restart the NFS service:

```
sudo service nfs-kernel-server restart
```

For NFS to work with the rest of the nodes, the Head node must always boot up before the slave nodes. If you encounter problems with a node not having access to the shared `/home/mpiuser` directory, restart the slave nodes so that the Head node is already on.

Ubuntu's firewall can also cause the nodes to have problems accessing the shared directory. To fix firewall problems we can set it so that it allows access from our network. Enter the following on the master node to do this:

```
sudo ufw allow from 10.0.0.0/24
```

7.2 Installing and Mounting NFS on Slave Nodes

For the slave nodes to be able to access and mount the shared directory, the following must be installed on each slave nodes:

```
sudo apt-get install nfs-common
```

Next we need to mount the directory to the slave nodes. Run the following on each of the slave nodes:

```
sudo mount headwolf:/home/mpiuser /home/mpiuser
```

To test if NFS worked, make a file in the `/home/mpiuser` directory on one node, and check if you can access it from another node.

Next we want to ensure that the directory is mounted automatically every time we turn on our nodes. do this by editing the `/etc/fstab` file:

```
sudo nano /etc/fstab
```

And add the following line:

```
headwolf:/home/mpiuser /home/mpiuser nfs
```

NFS should now be setup and each node should be sharing the same directory. Next we can go on the next step and install SSH for communication.

8 Setting Up SSH Password-less Communication

The cluster uses the Secure Shell (SSH) protocol to allow the Head node to control the rest of the nodes. Basically the head node can "login" to any of the nodes and run commands within that node. You can also manually take control of any node from any other node using SSH. Before you login SSH asks you for a pre-assigned password, meaning everytime you would run a job on the cluster, it would ask you to enter the password. We will setup passwordless SSH for more efficiency.

First install the SSH service into every node in the cluster by entering the following in every node:

```
sudo apt -get install openssh-server
```

On the Head node, generate an SSH key:

```
ssh -keygen -t rsa
```

Leave the field empty when asked for a password.

Next add this to the computer's list of authorized keys:

```
cd .ssh  
cat id_rsa.pub >> authorized_keys
```

Typically for the other nodes to have access, this key would need to be copied to each of the node's list of authorized keys. But since every node is sharing the same directory through NFS, the key is already in every node.

To test that SSH is working properly try logging into one node from a different node. For example on the Head node try to login to the fist node

```
ssh wolf01
```

You should now be logged into wolf01 from headwolf. Try shutting down wolf01 to verify functionality:

```
sudo poweroff
```

This should have shutdown the wolf01 node and logged you back into headwolf on the Head node.

If you are still prompted for a password when trying to login to a node using SSH, you need to setup a keychain by installing Keychain. Keychain is an ssh-agent that can store your key information in the computer's memory. Therefore you will only be prompted for a password once every time you reboot.

To install Keychain:

```
sudo apt-get install keychain
```

Next we want to edit the `/.bashrc` file and add the following:

```
if type keychain >/dev/null 2>/dev/null; then
    keychain --nogui -q id_rsa
    [ -f ~/.keychain/${HOSTNAME}-sh ] && . ~/.keychain/${HOSTNAME}-sh
    [ -f ~/.keychain/${HOSTNAME}-sh-gpg ] && . ~/.keychain/${HOSTNAME}-sh-gpg
fi
```

`/.bashrc` file is a shell script that automatically runs every time a new terminal is opened. This allows you to easily configure settings for your environment.

******* Bolding and lining are formatting errors please ignore.**

To make the changes exit and login once again or do:

```
source ~/.bashrc
```

9 Installing Compilers

10 Installing MPICH2

On the Head node, install MPICH2

```
sudo apt-get install mpich2
```

To verify that MPICH2 was installed correctly:

```
which mpiexec  
which mpirun
```

10.1 Getting Examples

The MPICH2 source package from the MPICH website contains a few examples that you can use to test your cluster. These examples are inside of the of the examples directory. You will need to compile these examples before using them.

in the /home directory of your head node, get the source package from the MPICH2 website:

```
sudo apt-get build-dep mpich2  
wget http://www.mpich.org/static/downloads/1.4.1/mpich2-1.4.1.tar.gz
```

Extract the source package and compile:

```
tar -xvzf mpich2-1.4.1.tar.gz  
cd mpich2-1.4.1.tar.gz  
./configure  
make
```

Navigate to the examples folder:

```
cd examples/
```

From here you can individually build these examples before using them. For example:

```
make cpi  
make child  
make ...
```

11 Testing and Benchmarks

12 Linux References

13 Other References

Links on Cluster

Links for basic networking

Links for Linux

Links for installing ubuntu

14 Appendix